



The University of Chicago Booth School of Business

Machine Learning

Winter 2020

HW #2

Jiyeon Chung

Arman Bhuiyan

Hikaru Sugimori

Deepak Putchakayala

1.1


The data-set includes historical car transaction data with a number of different data: price, year, color, region and more

1.2

```
carsdf = pd.read_csv("UsedCars.csv")
carsnp = carsdf.values
training_set, testing_set = train_test_split(carsnp, test_size=0.25, random_state=0)
```


1.3 Linear Regression

```
lm = LinearRegression()
lm.fit(X_train, y_train)
print(lm.intercept_)
plt.figure(figsize=(20,20))
plt.scatter(X_test, y_test, color='gray')
plt.xlabel("Mileage")
plt.ylabel("Price")
plt.plot(X_test, y_pred, color='red', label='linear Regression')
plt.gca().set_ylim(bottom=0)
plt.legend()
plt.show()
```



```
MSE = 124281428.121389
Intercept = 55438.601012
Coefficient = -0.337798
```

1.4 Polynomial Regression



Optimal polynomial degree = 7
MSE = 89876901.721073

```
mse_plot = [0] * 11  
degrees = list(range(11))
```

```
KF = KFold(n_splits=5, shuffle=True, random_state=0)  
folds = KF.split(mileage_train, price_train)
```

```
for fold in folds:
```

```
    train_index = list(fold[0])  
    cv_index = list(fold[1])  
    x_train = mileage_train[train_index, :]  
    y_train = price_train[train_index, :]  
    x_cv = mileage_train[cv_index, :]  
    y_cv = price_train[cv_index, :]
```

```
    for d in degrees:
```

```
        mse_temp = polynomial(d, x_cv, y_cv, x_train, y_train)  
        mse_plot[d] = mse_plot[d] + mse_temp
```

```
best_degree = min(mse_plot)
```


```
for i in range(11):
```

```
    if best_degree == mse_plot[i]:  
        best_degree = i  
    mse_plot[i] = mse_plot[i] / 5
```

```
plt.plot(degrees, mse_plot, color='blue')  
plt.show()
```

1.5

k-NN Regression



Optimal K = 460

MSE = 90161604.079432

```
mse_plot = [0] * 40
```

```
k = list(range(200, 600, 10))
```

```
KF = KFold(n_splits=5, shuffle=True, random_state=0)
```

```
folds = KF.split(mileage_train, price_train)
```

```
for fold in folds:
```

```
    train_index = list(fold[0])
```

```
    cv_index = list(fold[1])
```

```
    x_train = mileage_train[train_index, :]
```

```
    y_train = price_train[train_index, :]
```

```
    x_cv = mileage_train[cv_index, :]
```

```
    y_cv = price_train[cv_index, :]
```

```
for i in k:
```

```
    mse_temp = kNN(i, x_cv, y_cv, x_train, y_train)
```

```
    mse_plot[(i-200)//10] = mse_plot[(i-200)//10] + mse_temp
```

```
best_k = min(mse_plot)
```

```
for i in range(40):
```

```
    if best_k == mse_plot[i]:
```

```
        best_k = i
```

```
    mse_plot[i] = mse_plot[i] / 5
```

```
plt.plot(k, mse_plot, color='blue')
```

```
plt.show()
```

```
k-NN function
def knn(k, x_cv, y_cv, x_train, y_train):
    knn = KNeighborsRegressor(k, weights='uniform')
    y_predict = knn.fit(x_train, y_train).predict(x_cv)
    mse = mean_squared_error(y_cv, y_predict)

    return mse
```

Question 2

Principal component analysis (PCA) is a technique used to emphasize variation and bring out strong patterns in a dataset. PCA is essentially just a coordinate transformation. The original data are plotted on an X-axis and a Y-axis. For two-dimensional data, PCA seeks to rotate these two axes so that the new axis X' lies along the direction of maximum variation in the data. For this problem, we will generate and analyze cluster in the dataset using this technique.

We merged the wine deals and wine transactions datasets since they shared a column, created a pivot table to show customer response successes per offer, sliced the matrix to isolate only the binary indicator column, counted the number of people who ended up in each cluster and plotted them. We also transformed the two dimensional dataset via PCA, created a new df with name, cluster membership, and coordinates, merged that dataframe, and created a scatterplot. For more details on how this was done, please see the code at the end.

```
[In [9]: averages = review_cluster(merged_df, 4)
```



```
[In [10]: averages
```

```
Out[10]:
```

	Discount (%)	Minimum Qty	cluster
True	60.639175	75.030928	1.0
True	64.652174	79.617391	2.0
True	47.387097	51.483871	3.0
True	55.830508	6.000000	4.0

The screenshot above describes the clusters we got in terms of discount % and minimum quantity. Cluster 1 and Cluster 2 are very similar in terms of their quantity and discount, so it's likely the varietal or origins that are different. Cluster 4 takes the least discount and the least quantity, so if there are also fewer of them, these are people that don't buy much. Cluster 3 is somewhere in the middle.

Visualisations of our clusters are below. Looks like most people belong to cluster 1 and 3.



8]: min_quantity discount

is_4

False	51.512727	59.287273
True	97.102041	60.571429

```

(Varietal
Champagne          57
Cabernet Sauvignon 32
Malbec             15
Chardonnay         14
Merlot             13
Prosecco           11
Pinot Noir         7
Espumante          6
Pinot Grigio       2
Name: Varietal, dtype: int64, Origin
France            78
Chile             24
Italy            12
Oregon           10
California        9
New Zealand       7
Australia         6
South Africa      6
Germany           5
Name: Origin, dtype: int64)

```

First cluster liked a variety of origins and varieties, driven by France and champagne

```

(Varietal
Espumante          25
Malbec             12
Pinot Grigio       9
Prosecco           6
Merlot             4
Champagne          1
Name: Varietal, dtype: int64, Origin
France            21
South Africa      16
Oregon            9
Chile             5
Australia         4
California        1
New Zealand       1
Name: Origin, dtype: int64)

```

The second cluster mostly liked espumante and malbec from France and South Africa.


```
(Varietal
Champagne      20
Prosecco       9
Espumante      5
Name: Varietal, dtype: int64, Origin
France         11
Chile          6
New Zealand    4
Australia      3
Germany        3
Oregon         3
California     2
South Africa   2
Name: Origin, dtype: int64)
```

The third cluster liked even fewer things, but mostly champagne (see above).

```
(Varietal
Pinot Noir     37
Prosecco       2
Champagne      1
Chardonnay     1
Malbec         1
Merlot         1
Name: Varietal, dtype: int64, Origin
Australia      13
Italy          12
France         7
Germany        7
California     2
New Zealand    1
South Africa   1
Name: Origin, dtype: int64)
```

The fourth cluster basically only drinks pinot from Australia or Italy

The code for #2 is below.

```
def collect_data():
    wine_deals = pd.read_csv("Wine_deals.csv")
    wine_transactions = pd.read_csv("Wine_transactions.csv")
    wine_transactions['n'] = 1
    df = pd.merge(wine_deals, wine_transactions)
    df = df[['Offer #', 'Campaign', 'Customer Last Name', 'Origin', 'Varietal', 'Discount (%)', 'Minimum
Qty', 'Past Peak', 'n']]
    return df, wine_transactions, wine_deals

def pivot(df):
    # Pivot table to show customer response successes per offer
    matrix = df.pivot_table(index=['Customer Last Name'], columns=['Offer #'], values='n')
    matrix = matrix.fillna(0).reset_index()
```

```

x_columns = matrix.columns[1:]
cluster = KMeans(n_clusters=5)
matrix['cluster'] = cluster.fit_predict(matrix[matrix.columns[2:]])
matrix.cluster.value_counts()
pca = PCA(n_components=2)
matrix['x'] = pca.fit_transform(matrix[x_columns])[:,0]
matrix['y'] = pca.fit_transform(matrix[x_columns])[:,1]
matrix = matrix.reset_index()
return matrix

def plot_clusters(matrix):
    sns.countplot(x='cluster', data=matrix, order=matrix['cluster'].value_counts().index)
    plt.show()

def merge_cluster_dataframe(matrix, wine_transactions, wine_deals):
    customer_clusters = matrix[['Customer Last Name', 'cluster', 'x', 'y']]
    df = pd.merge(wine_transactions, customer_clusters)
    df = pd.merge(wine_deals, df)
    return df

def plot_merged_dataframe(df):
    sns.lmplot('x', 'y', data=df, hue='cluster', fit_reg=False, palette='deep', legend=True, size=4)
    plt.show()

def review_cluster(df, cluster_number):
    averages = pd.DataFrame()
    for x in range(0, cluster_number):
        df[x] = df.cluster == x
        subset = df.groupby(x)[['Discount (%)', 'Minimum Qty']].mean()
        subset['cluster'] = x+1
        averages = averages.append(subset.loc[True])
    return averages

df, wine_transactions, wine_deals = collect_data()
matrix = pivot(df)
merged_df = merge_cluster_dataframe(matrix, wine_transactions, wine_deals)
plot_merged_dataframe(merged_df)
averages = review_cluster(merged_df, 4)

```

